



kubernetes

Advanced Concepts

Kubernetes



Course Objectives

In this module, you will learn:

- Commands & Arguments
- Environment variables
- Config Maps & Secrets
- Health Checks
- Jobs/ Cron Jobs
- Application Logs
- Daemon Sets / Stateful Sets



kubernetes

Commands & Arguments

Advanced Concepts



Dockerfile vs Pod YAML

```
FROM Ubuntu
```

```
ENTRYPOINT ["sleep"]
```

```
CMD ["5"]
```

pod-definition.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper-pod
spec:
  containers:
    - name: ubuntu-sleeper
      image: ubuntu-sleeper
      args: ["10"]
```

```
▶ kubectl create -f pod-definition.yml
```

Dockerfile vs Pod YAML

FROM Ubuntu

ENTRYPOINT ["sleep"]

CMD ["5"]

pod-definition.yml

apiVersion: v1

kind: Pod

metadata:

name: ubuntu-sleeper-pod

spec:

containers:

- name: ubuntu-sleeper

image: ubuntu-sleeper

command: ["sleep2.0"]

args: ["10"]



kubernetes

Environment variables in Applications

Advanced Concepts



ENV variables in Kubernetes

```
▶ docker run -e APP_COLOR=pink simple-webapp-color
```

pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp-color
spec:
  containers:
  - name: simple-webapp-color
    image: simple-webapp-color
    ports:
    - containerPort: 8080
    env:
    - name: APP_COLOR
      value: pink
```

ENV Value Types

env:

- **name:** APP_COLOR
value: pink



env:

- **name:** APP_COLOR
valueFrom:
 configMapKeyRef:



env:

- **name:** APP_COLOR
valueFrom:
 secretKeyRef:





kubernetes

Configuring Config Maps

Advanced Concepts



ConfigMap

- In programming, we use env files or separate configuration files to store settings, configurations, or variables that are required to execute the program.
- In Kubernetes, we can use ConfigMaps to achieve the same functionality.
- A ConfigMap is a Kubernetes API object that can be used to store data as key-value pairs. Kubernetes pods can use the created ConfigMaps as a:
 - Configuration file
 - Environment variable
 - Command-line argument


ConfigMap in Pods

pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: simple-webapp-color
  labels:
    name: simple-webapp-color
spec:
  containers:
  - name: simple-webapp-color
    image: simple-webapp-color
    ports:
      - containerPort: 8080
    envFrom:
      - configMapRef:
          name: app-config
```

config-map.yaml

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  APP_COLOR: blue
  APP_MODE: prod
```



 `kubectl create -f pod-definition.yaml`

View ConfigMaps

```
▶ kubectl get configmaps
```

NAME	DATA	AGE
app-config	2	3s

```
▶ kubectl describe configmaps
```

```
Name:          app-config
Namespace:     default
Labels:        <none>
Annotations:   <none>
```

```
Data
```

```
====
```

```
APP_COLOR:
```

```
----
```

```
blue
```

```
APP_MODE:
```

```
----
```

```
prod
```

```
Events:  <none>
```



kubernetes

Configuring Secrets

Advanced Concepts



Kubernetes Secrets

- A secret is an object that contains a small amount of sensitive data
 - Such as a password, a token, or a key
- Putting this information in a secret is safer and more flexible than putting it verbatim in a pod definition or in a Docker image
 - A secret is only sent to a node if a pod on that node requires it
 - Not written to disk—stored in a tmpfs on the nodes
 - Deleted once the pod that depends on it is deleted
 - One pod does not have access to the secrets of another pod
 - Secrets are encrypted at the storage layer in etcd

Creating Kubernetes Secrets

```
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  username: YWRtaW4=
  password: MWYyZDFlMmU2N2Rm
```

Not encrypted, base64 encoded

```
$ kubectl get secret mysecret -o yaml
```

To get the secret



kubernetes

Health Checkers

Application Lifecycle Management



Health Checkers

There are two types of Kubernetes health checks: Liveness and Readiness probes

- If a Liveness probe fails, the pod is deleted and re-created
- If a Readiness probe fails, requests are not sent to the pod until it is back up and running
 - But it is not removed and recreated

POD Conditions

PodScheduled	TRUE	FALSE
Initialized	TRUE	FALSE
ContainersReady	TRUE	FALSE
Ready	TRUE	FALSE

```
kubectl describe pod

Name:          nginx-65899c769f-9lwzh
Namespace:     default
Node:          kubenode2/192.168.1.103
Start Time:    Wed, 08 Aug 2018 22:57:39 -0400
Labels:        pod-template-hash=2145573259
               run=nginx
Annotations:    <none>
Status:        Running
IP:            10.244.2.222
Controlled By: ReplicaSet/nginx-65899c769f
Containers:
  nginx:
    Image:          nginx
    Image ID:       docker-
pullable://nginx@sha256:d85914d547a6c92faa39ce7058bd7529baa
cab7e0cd4255442b04577c4d1f424
    Port:           <none>
    Host Port:      <none>
    State:          Running
      Started:      Wed, 08 Aug 2018 22:57:55 -0400
      Ready:        True
default-token-hxr6t (ro)
Conditions:
  Type           Status
  Initialized     True
  Ready           True
  PodScheduled    True
```

Readiness/ Liveness Probes



HTTP Test - /api/ready



TCP Test - 3306



Exec Command

Liveness and Readiness Probe Configuration

```
apiVersion: apps/v1
kind: Deployment
*** CODE OMITTED ***
spec:
  containers:
    *** CODE OMITTED ***
    readinessProbe:
      httpGet:
        path: /health
        port: 8080
        initialDelaySeconds: 30
        periodSeconds: 30
    livenessProbe:
      httpGet:
        path: /health
        port: 8080
        initialDelaySeconds: 15
        periodSeconds: 15
```

Added to the Containers section of the Deployment



kubernetes

Jobs/ Cron Jobs

Application Lifecycle Management



Job

- A Kubernetes Job is a workload controller that represents a finite task.
- Jobs differ from other controller objects in that Jobs manage the task as it runs to completion, rather than managing an ongoing desired state (such as the total number of running Pods).
- When a specified number of pods reach completion, the Job is said to have successfully completed.

Job – Use cases

- The best use case for Kubernetes jobs are,
- **Batch processing:** Let's say you want to run a batch task once a day or during a specific schedule. It could be something like reading files from storage or a database and feed them to a service to process the files.
- **Operations/ad-hoc tasks:** Let's say you want to run a script/code which runs a database cleanup activity or to even backup a kubernetes cluster itself.

Restart Policy – Never vs Always

pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: math-pod
spec:
  containers:
  - name: math-add
    image: ubuntu
    command: ['expr', '3', '+', '2']
    restartPolicy: Never
```

pod-definition.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: math-pod
spec:
  containers:
  - name: math-add
    image: ubuntu
    command: ['expr', '3', '+', '2']
    restartPolicy: Always
```


Create, View & Delete

job-definition.yaml

```
apiVersion: batch/v1
kind: Job
metadata:
  name: math-add-job
spec:
  template:
    spec:
      containers:
        - name: math-add
          image: ubuntu
          command: ['expr', '3', '+', '2']

      restartPolicy: Never
```

```
▶ kubectl create -f job-definition.yaml
```

```
▶ kubectl get jobs
```

NAME	DESIRED	SUCCESSFUL	AGE
math-add-job	1	1	38s

```
▶ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
math-add-job-187pn	0/1	Completed	0	2m

```
▶ kubectl logs math-add-job-1d87pn
```

```
5
```

```
▶ kubectl delete job math-add-job
```

```
job.batch "math-add-job" deleted
```

Multiple Job Pods and Parallelism

- When a job is deployed you can make it run on multiple pods with parallelism.
- For example, in a job if you want to run 6 pods and run 2 pods in parallel, you need to add the following two parameters to your job manifest.

completions: 3

parallelism: 3

- The job will run 2 pods in parallel for 3 times to achieve 6 completions.
- Here is the manifest file with those parameters.

Multiple Job Pods and Parallelism

job-definition.yaml

```
apiVersion: batch/v1
kind: Job
metadata:
  name: random-error-job
spec:
  completions: 3
  parallelism: 3
  template:
    spec:
      containers:
      - name: random-error
        image: random-error
      restartPolicy: Never
```

```
kubectl create -f job-definition.yaml
```

```
kubectl get jobs
```

NAME	DESIRED	SUCCESSFUL	AGE
random-error-job	3	2	38s

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS
random-exit-job-ktmtt	0/1	Completed	0
random-exit-job-sdsrf	0/1	Error	0
random-exit-job-wwqbn	0/1	Completed	0
random-exit-job-fkhfn	0/1	Error	0
random-exit-job-fvf5t	0/1	Error	0
random-exit-job-nmghp	0/1	Completed	0



Jobs



kubernetes

Cron Jobs

Application Lifecycle Management



Cronjob

- What if you want to run a batch job on specific schedules, for example, every 2 hours. You can create a Kubernetes cronjob with a cron expression.
 - The job will automatically kick in as per the schedule you mention in the job.
- A CronJob is basically a Kubernetes Job with time-based scheduling and some specific parameters to handle failure.

CronJob

job-definition.yaml

```
apiVersion: batch/v1
kind: Job
metadata:
  name: reporting-job
spec:
  completions: 3
  parallelism: 3
  template:
    spec:
      containers:
        - name: reporting-tool
          image: reporting-tool
      restartPolicy: Never
```

cron-job-definition.yaml

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: reporting-cron-job
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
```

CronJob

job-definition.yaml

```
apiVersion: batch/v1
kind: Job
metadata:
  name: reporting-job
```

cron-job-definition.yaml

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: reporting-cron-job
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      completions: 3
      parallelism: 3
      template:
        spec:
          containers:
            - name: reporting-tool
              image: reporting-tool

  restartPolicy: Never
```



kubernetes

Managing Application Logs

Monitoring & Logging



Application Logs

- Application logs are the logs from the applications that run on Kubernetes.
- The data stored in these logs consists of the information that your applications output as they run.
- Typically, this data is written to stdout inside the container where the application runs.

Viewing Application Logs

- There are two main ways to interact with application log data.
- The first is to run a command like

`kubectl logs pod-name`

where “pod-name” is the name of the pod that hosts the application whose logs you want to access.

```
kubectl create -f event-simulator.yaml
```

```
kubectl logs -f event-simulator-pod
```

```
2018-10-06 15:57:15,937 - root - INFO - USER1 logged in
2018-10-06 15:57:16,943 - root - INFO - USER2 logged out
2018-10-06 15:57:17,944 - root - INFO - USER2 is viewing page2
2018-10-06 15:57:18,951 - root - INFO - USER3 is viewing page3
2018-10-06 15:57:19,954 - root - INFO - USER4 is viewing page1
2018-10-06 15:57:20,955 - root - INFO - USER2 logged out
2018-10-06 15:57:21,956 - root - INFO - USER1 logged in
2018-10-06 15:57:22,957 - root - INFO - USER3 is viewing page2
2018-10-06 15:57:23,959 - root - INFO - USER1 logged out
2018-10-06 15:57:24,959 - root - INFO - USER2 is viewing page2
2018-10-06 15:57:25,961 - root - INFO - USER1 logged in
2018-10-06 15:57:26,965 - root - INFO - USER4 is viewing page3
2018-10-06 15:57:27,965 - root - INFO - USER4 is viewing page3
2018-10-06 15:57:28,967 - root - INFO - USER2 is viewing page1
2018-10-06 15:57:29,967 - root - INFO - USER3 logged out
2018-10-06 15:57:30,972 - root - INFO - USER1 is viewing page2
```

event-simulator.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: event-simulator-pod
spec:
  containers:
  - name: event-simulator
    image: event-simulator
```

Storing & Analyzing Application Logs

- Suppose you want to store logs persistently and analyze them systematically.
- In that case, you're better served by using an external logging tool like IBM Log Analysis with LogDNA to collect and interpret the logs.
- The easiest way to go about this is to run a so-called sidecar container, which runs alongside the application, collects its logs, and makes them available to an external logging tool.



kubernetes

Daemon Sets & Stateful Sets

Monitoring & Logging



DaemonSet - When Your Pods Must Be Everywhere

- A DaemonSet is a controller that ensures that the pod runs on all the nodes of the cluster.
 - If a node is added/removed from a cluster, DaemonSet automatically adds/deletes the pod.
- Some typical use cases of a DaemonSet is to run cluster level applications like:
 - Monitoring Services
 - Logs Collection Daemon
 - Cluster services like kube-proxy
 - Security solutions like antivirus, intrusion detection, and image scanning

StatefulSet - When Your Pods Are Unique

- A StatefulSet is a controller that helps you deploy and scale groups of Kubernetes pods.
- When using Kubernetes, most of the time you don't care how your pods are scheduled, but sometimes you care that pods are deployed in order, that they have a persistent storage volume, or that they have a unique, stable network identifier across restarts and reschedules.
- In those cases, StatefulSets can help you accomplish your objective.
- To scale *stateful systems*, you need a **StatefulSet**.

StatefulSet - When Your Pods Are Unique

- Each pod created by the StatefulSet has an ordinal value (0 through # replicas - 1) and a stable network ID (which is statefulsetname-ordinal) assigned to it.
- You can also create a **VolumeClaimTemplate** in the manifest file that will create a persistent volume for each pod.
- When pods are deployed by a StatefulSet, they will go in order from 0 to the final pod and require that each pod is Running and Ready before creating the next pod.

StatefulSet - When to use?

- Some examples of reasons you'd use a StatefulSet include:
 - A Redis pod that has access to a volume, but you want it to maintain access to the same volume even if it is redeployed or restarted
 - A Cassandra cluster and have each node maintain access to its data
 - A webapp that needs to communicate with its replicas using known predefined network identifier