



kubernetes

# Networking Fundamentals

Kubernetes




# Course Objectives

In this module, you will learn:

- Networking at the Pod level
- Networking at the Cluster level

# Pods with IP Addresses

- IP address are assigned to a Pod
- Containers in the same pod communicate using localhost



# Internet-to-Service Networking

- We need to be able to access our deployments from outside the cluster:
- When setting up external access, there are two techniques you need to use — egress and ingress.
- These are policies that you can set up with either whitelisting or blacklisting to control traffic into and out of your network.

# Internet-to-Service Networking

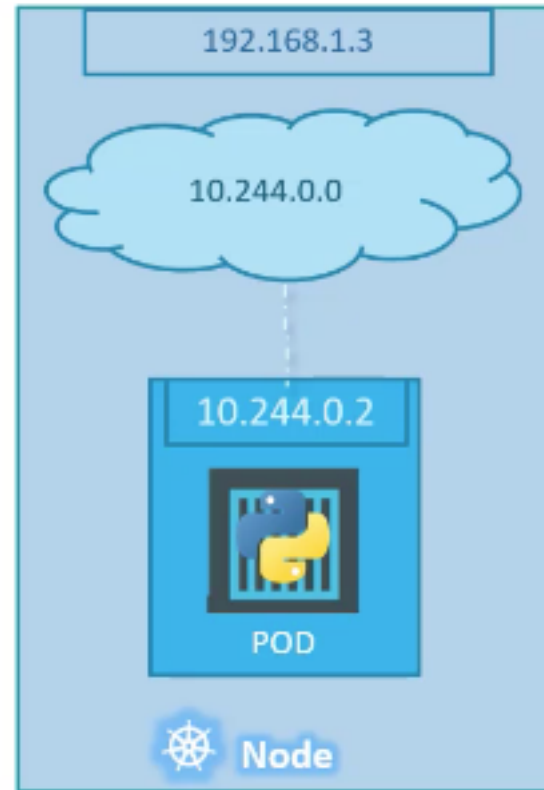
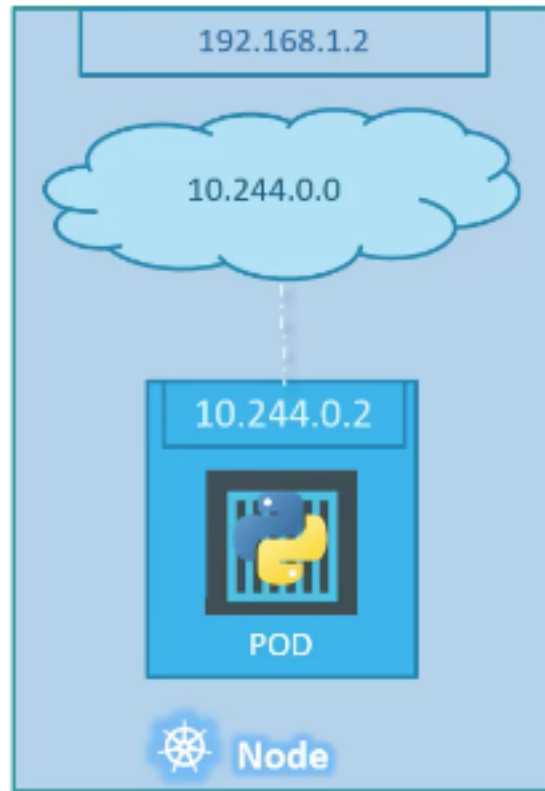
## ■ Egress:

- Egress is a process that routes traffic from your node to an outside connection.
- It is often accomplished via an Internet gateway attached to your virtual private cloud (VPC).
- This gateway uses network address translation ([NAT](#)) to map IPs between your users and the machine your node is hosted on.
- It cannot, however, map to the individual pods on your node.
- For this step, Kubernetes uses IP tables and cluster-IPs to finalize communications.

## ■ Ingress:


- Ingress is an opposite process to egress and involves communications from external clients to your Kubernetes services.
- It operates as a set of rules defining what connections are allowed and which are blocked from communicating with your services.

# Cluster Networking



# Cluster Networking

- All containers/PODs can communicate to one another without NAT
- All nodes can communicate with all containers and vice versa without NAT



# Networking Tools





# Cluster Networking

